Database & SQL

Induction Training

Do The Math

Chicago, IL
Bangalore, India
www.mu-sigma.com

January 16, 2013

Welcome to SQL & Database Training

So… what is SQL and why do we need it?

SQL stands for:
Scarcely Qualified Language

Just kidding!

So… what is SQL and why do we need it?

SQL stands for:
Structured Query Language

That's all you will learn for now!

Kidding!!... But we will hold off SQL for sometime

Let us spend two minutes on this…

What are your expectations from this training?

Before we dive into SQL, let's talk about DATA…

Sorry… when we say DATA,
we mean **information**

Let’s try and understand where and how data is created

A Retail Supermarket

A Customer

What kind of information is captured?

A retail store captures a lot of information regarding a customer's shopping behavior

Who is coming to my store?

What products are being bought?

What frequently are my products purchased?

Let’s list out some metrics or data a retail store might collect

1 Customer Information

- Customer Name
- Customer Address
- Age
- Telephone

2 Product Information

- Product Name
- Product Class
- Price
- Discount

2 Transaction Information

- Date
- Customer Name
- Product Name
- Cost

Now, consider purchases of a single customer

Dr. John Watson - the customer

- Bread • 1 lb. $3.02
- Eggs • 2.5 Doz., $0.30
- Milk • 1 Gal. $2.49
- Apples • 2 lbs., $2.34

How do we store this information?

What if we use Excel to store this information?

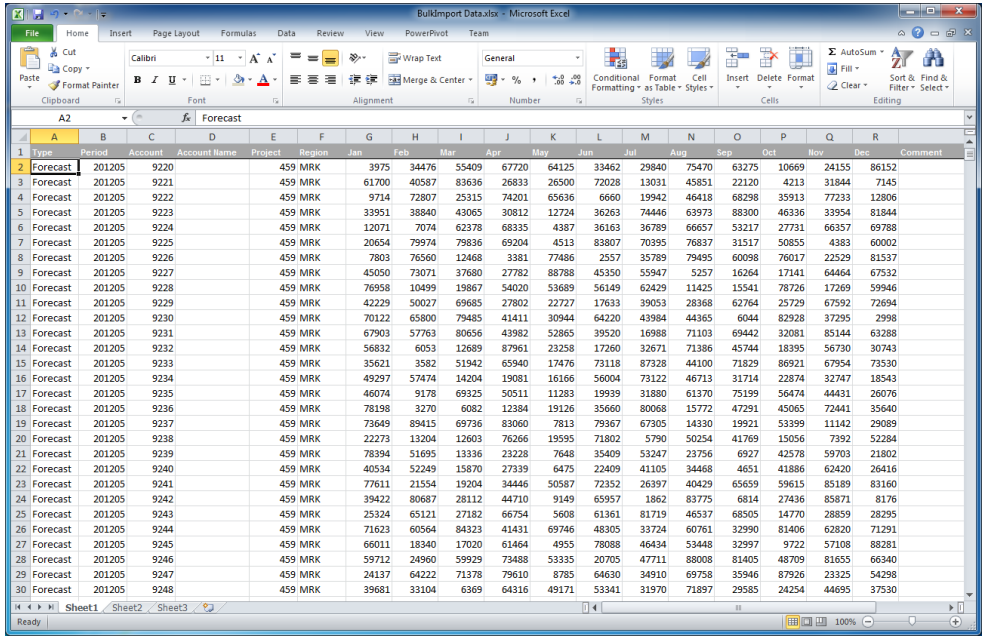

Excel Spreadsheet

- What if 100 people want to modify the file?
- What if the number of transactions cross a million?
- How do you manage a million pieces of information?
- **An Excel document is clearly not the right choice for storing information in a transactional environment**
- **What we need is a file system which can**
  - **Hold multiple tables of information**
  - **Do not run out of space**
  - **Allow multiple people to access and modify the information simultaneously**

A database allows multiple users to interact with the same piece of information simultaneously

InterBase Firebird Panorama

SQLBase **SQL Server** Sybase

Filemaker **RDBMS** Access

FoxPro **Oracle** **DB2** SQLite

PostGreSQL

MaxDB **MySQL** Teradata

- More on **RDBMS** in the later slides …

How should we store every transaction in the database?

Store every transaction as a single row

Date	First Name	Last Name	Bread	Eggs	Milk	Apple
3-Mar-12	John	Watson	$3.02	$0.30	$2.49	$2.34

Advantages

- I can find everything about a purchase by looking at single row

Disadvantages

- Not feasible as the number of columns is undefined (if a customer buys 20 items)

How should we store every transaction in the database?

Store every item in a transaction as a single row

Date	First Name	Last Name	Product	Quantity	Price
3-Mar-12	John	Watson	Bread	1 lb.	$3.02
3-Mar-12	John	Watson	Eggs	2.4 doz.	$0.30
3-Mar-12	John	Watson	Milk	1 gal.	$2.49
3-Mar-12	John	Watson	Apple	2 lbs.	$2.34

Advantages

- I don't have to worry about how many items a customer is buying

Disadvantages

- There is duplication of information across the rows

That brings us to the subject of a relational database system

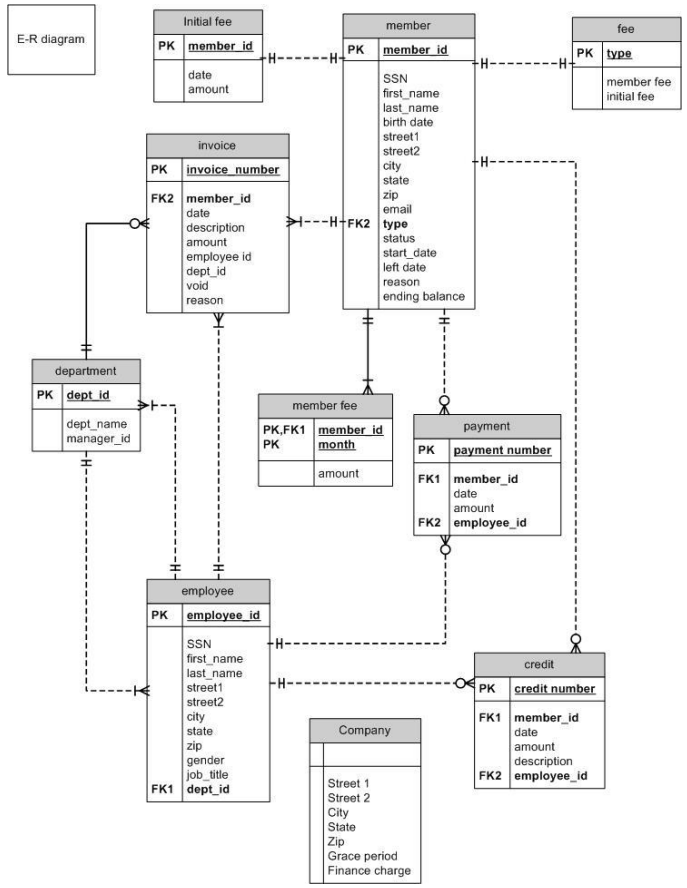

- A database is just a location to store and retrieve data
- A relational database is one which treats all of its data as a collection of relations
- Each table in a relational database has unique key (also referred to as the primary key)
- This primary key can also be present in another table as a foreign key and in process creates a relation between two tables

To create a relational database, we need to understand the granularity or level of each table

StudentID	FirstName	LastName	CourseID
S002001	Jack	Black	C004
S002002	Lucy	Liu	C001
S002003	Dustin	Hoffman	C002
S002004	Seth	Rogen	C004

Foreign Key

Relationship

Primary Key

A primary key uniquely identifies each record in a table

CourseID	CourseName
C001	History
C002	Computing
C003	Biology
C004	English

A simple relational database

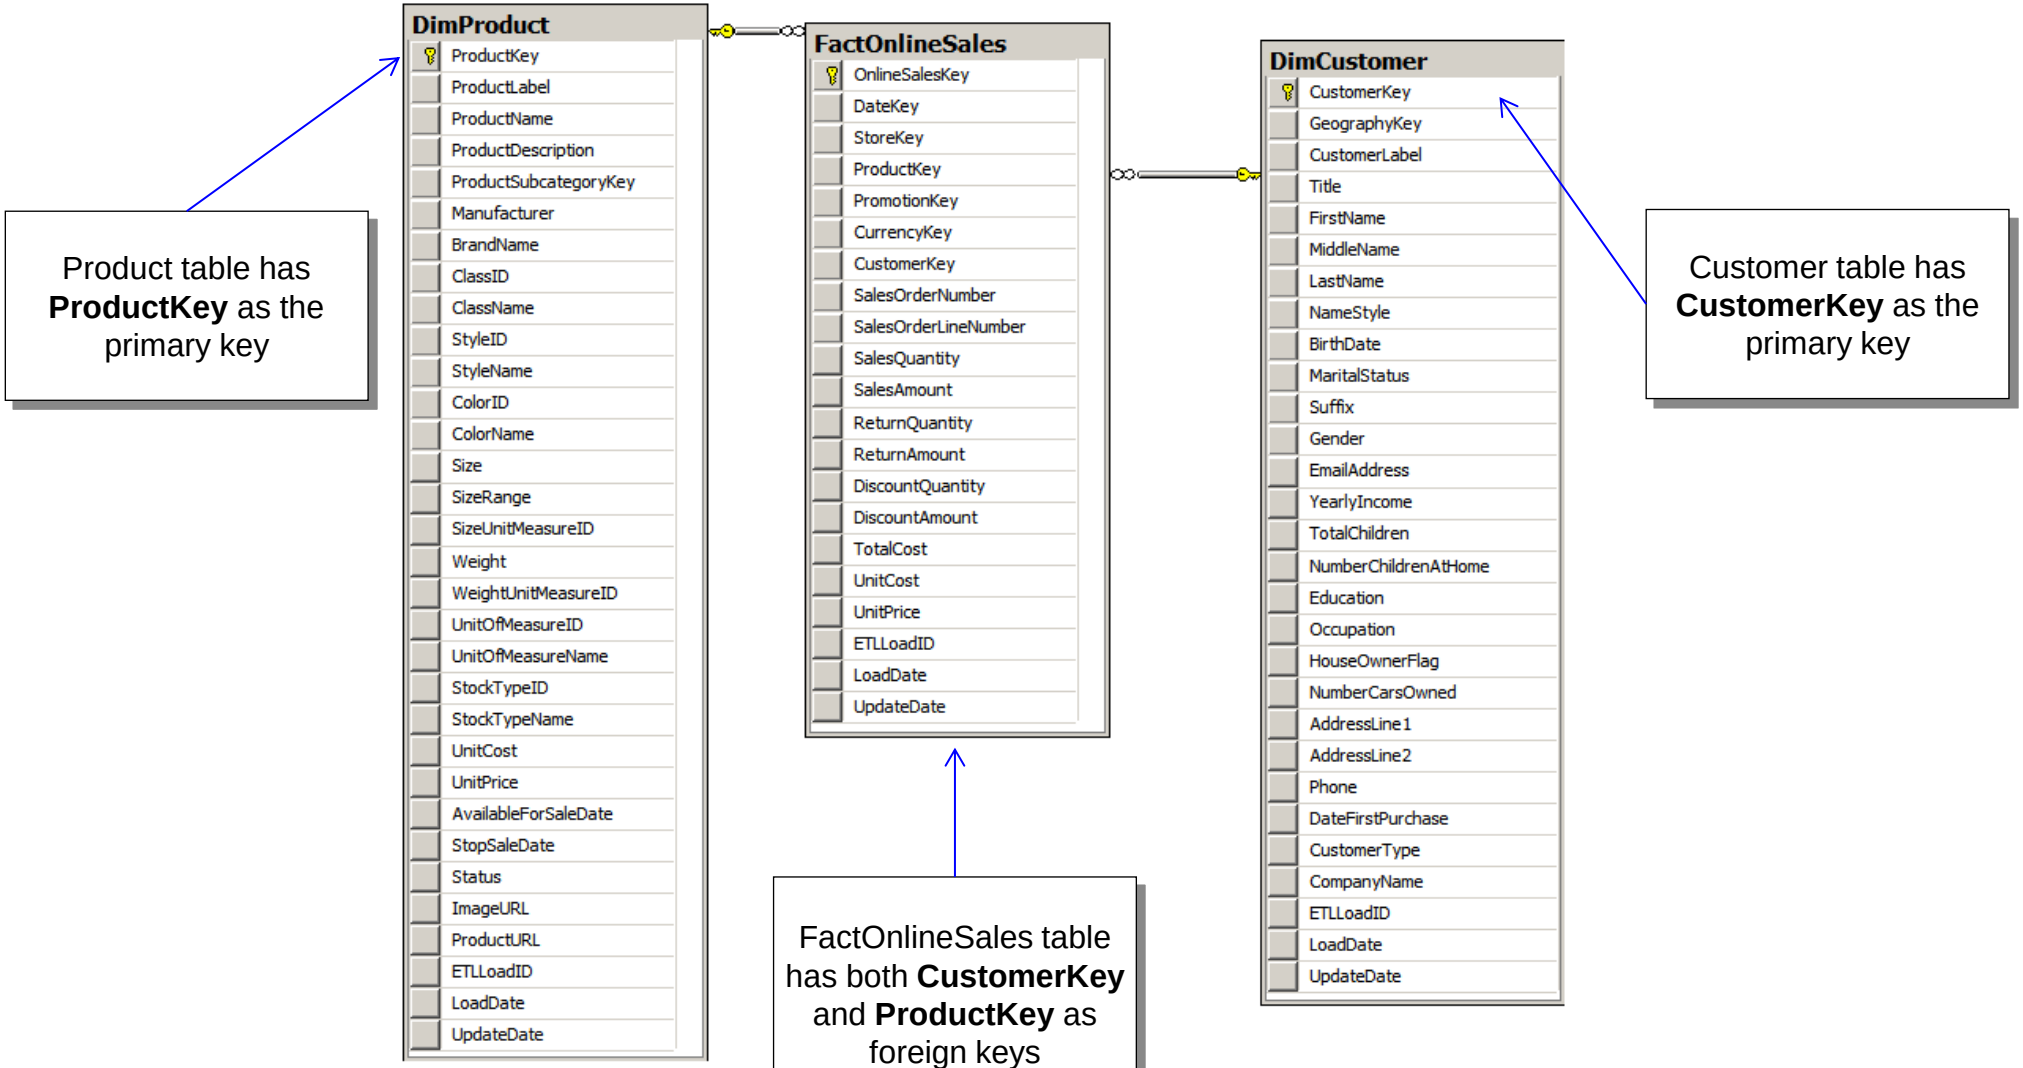

A database management system is not just a place to dump data

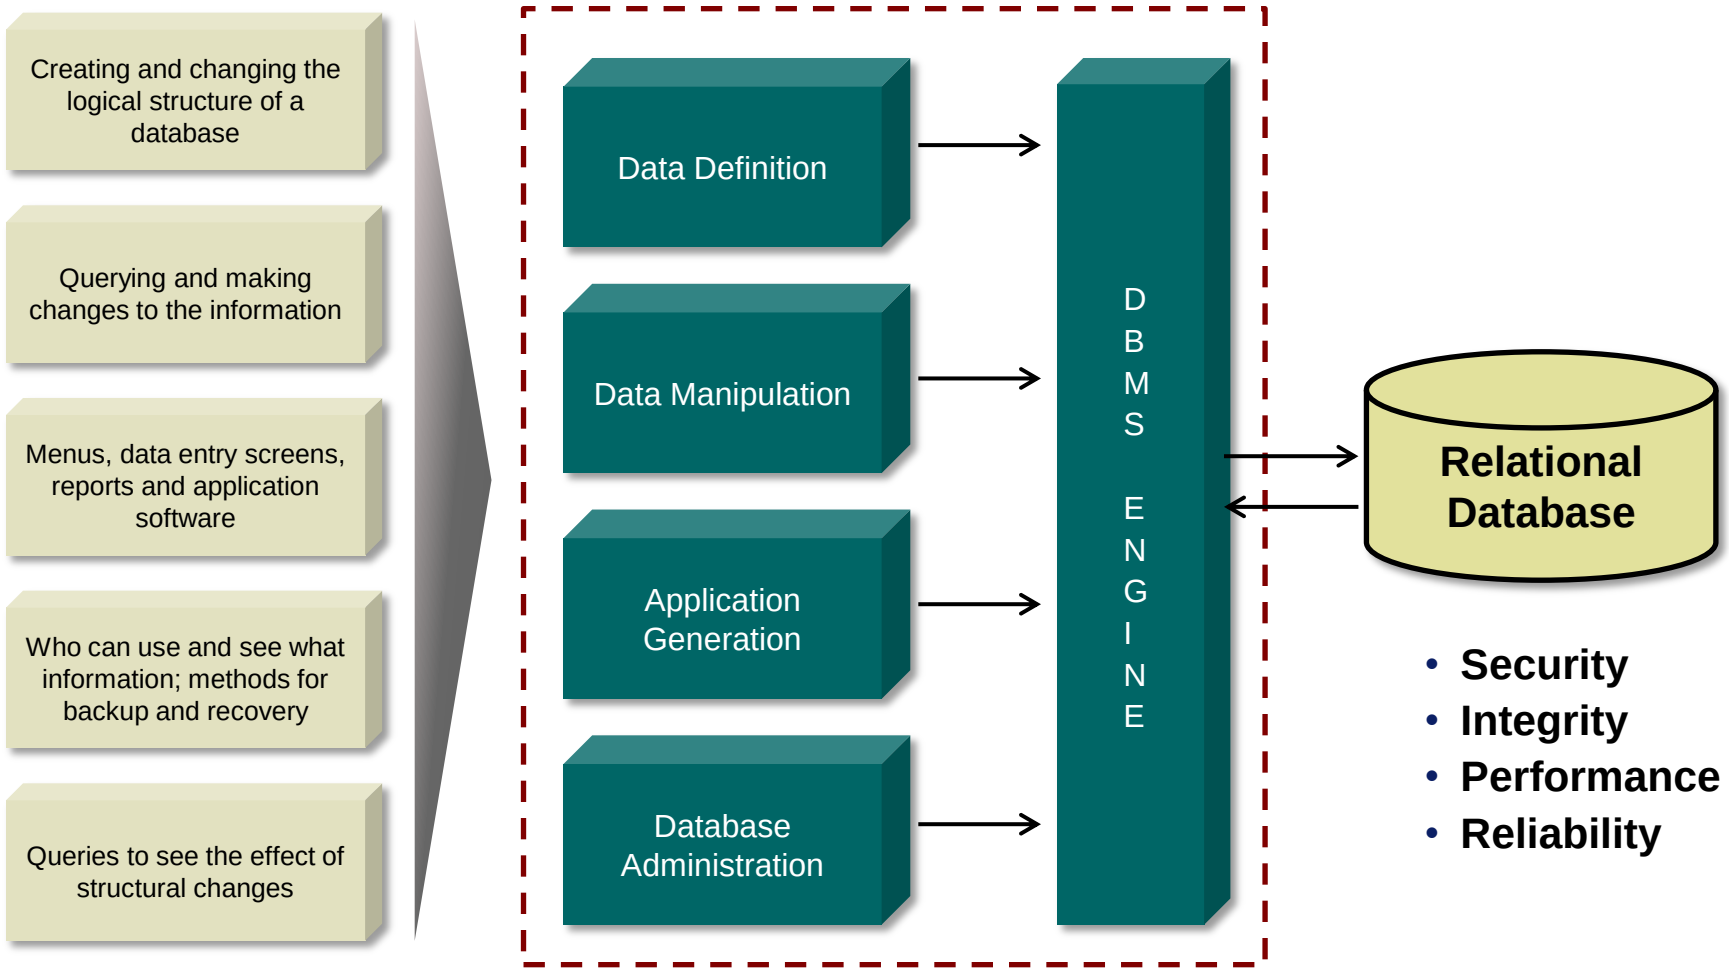

Coming back to what is SQL and why do we need it?

- SQL is the standard language for communicating with relational database management systems
- SQL statements are used to perform tasks such as update data on a database, or retrieve data from a database
- Although most database systems use SQL, most of them also have their own additional proprietary extensions that are usually only used on their system
- However the standard SQL commands can be used to accomplish almost everything one needs

For our learning purpose we will be using the Microsoft SQL Server 2008 Express edition

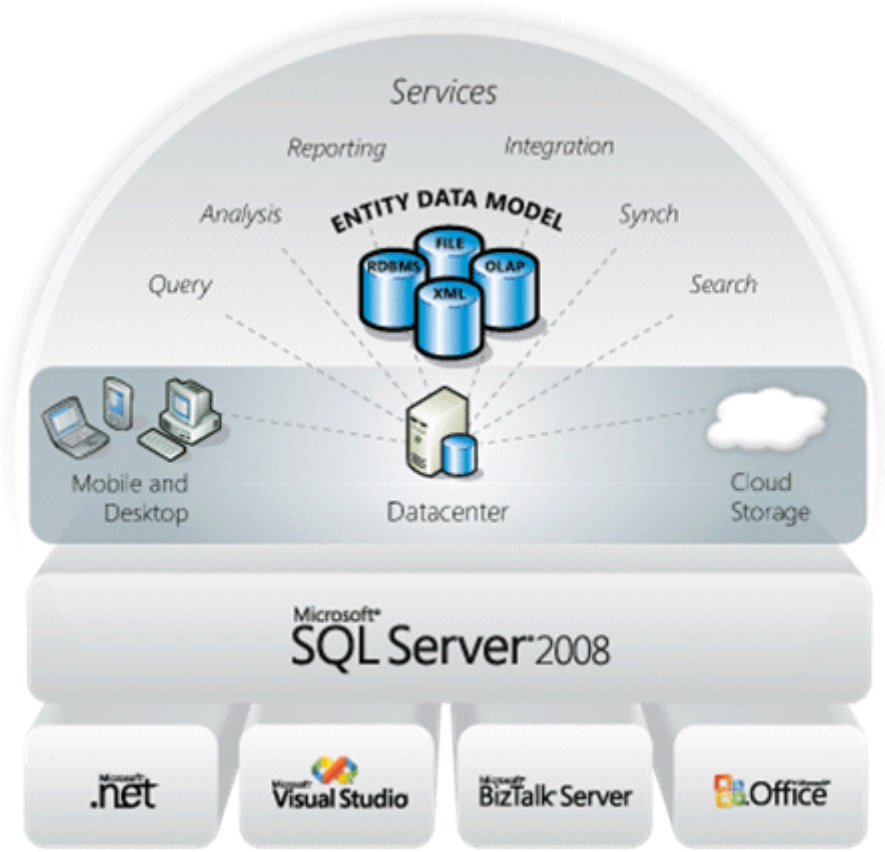

- SQL Server 2008 R2 Express edition is a free version of SQL Server 2008 R2
- Almost everything that can be done in the commercial version is supported in the express edition
- Database size is restricted to 10 GB in size
- Uses only a single processor
- Uses only 1 GB of internal memory

Start up SQL Server Management Studio from your Start menu

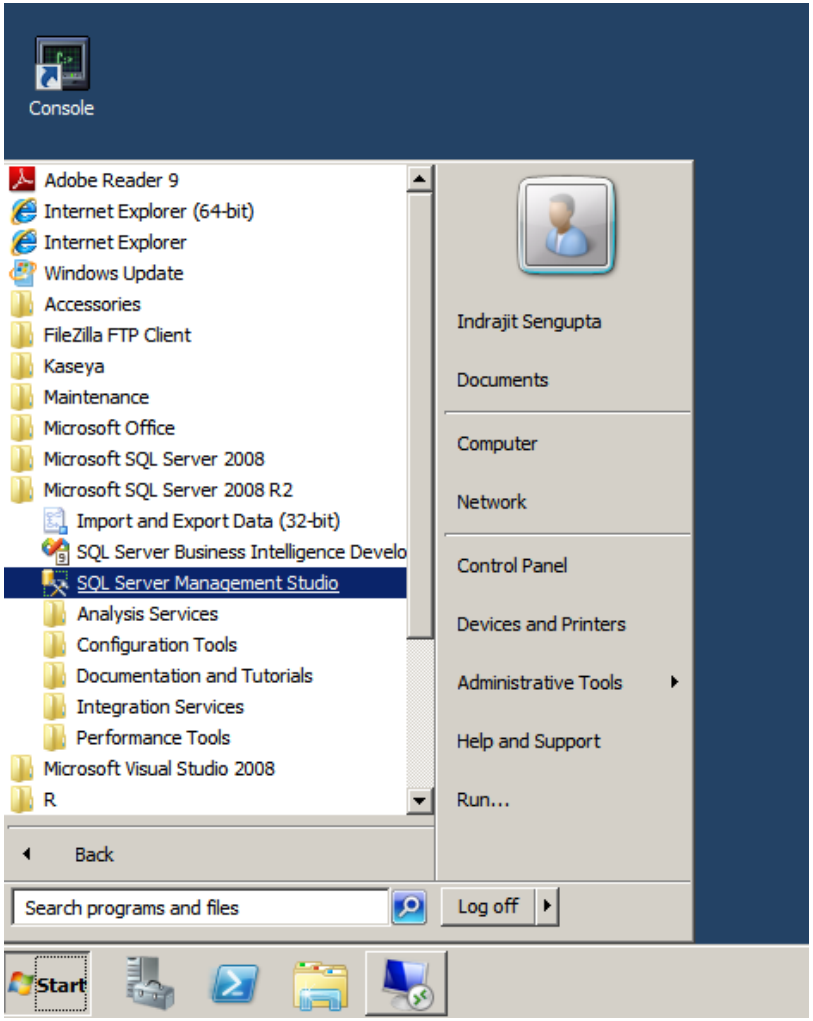

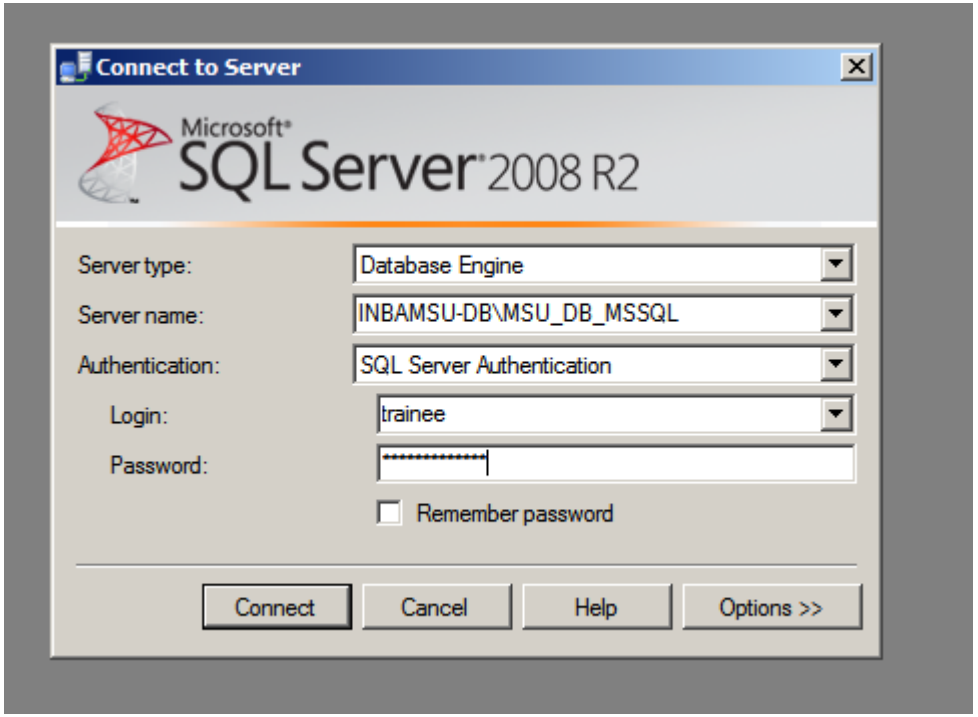

- Enter the server name: **INBAMSU-DB\MSU_DB_MSSQL**
- Enter user id: **trainee**
- Enter password: **password@1234**
- Remember to choose **SQL Server Authentication**

The Object Explorer window will list out all the databases

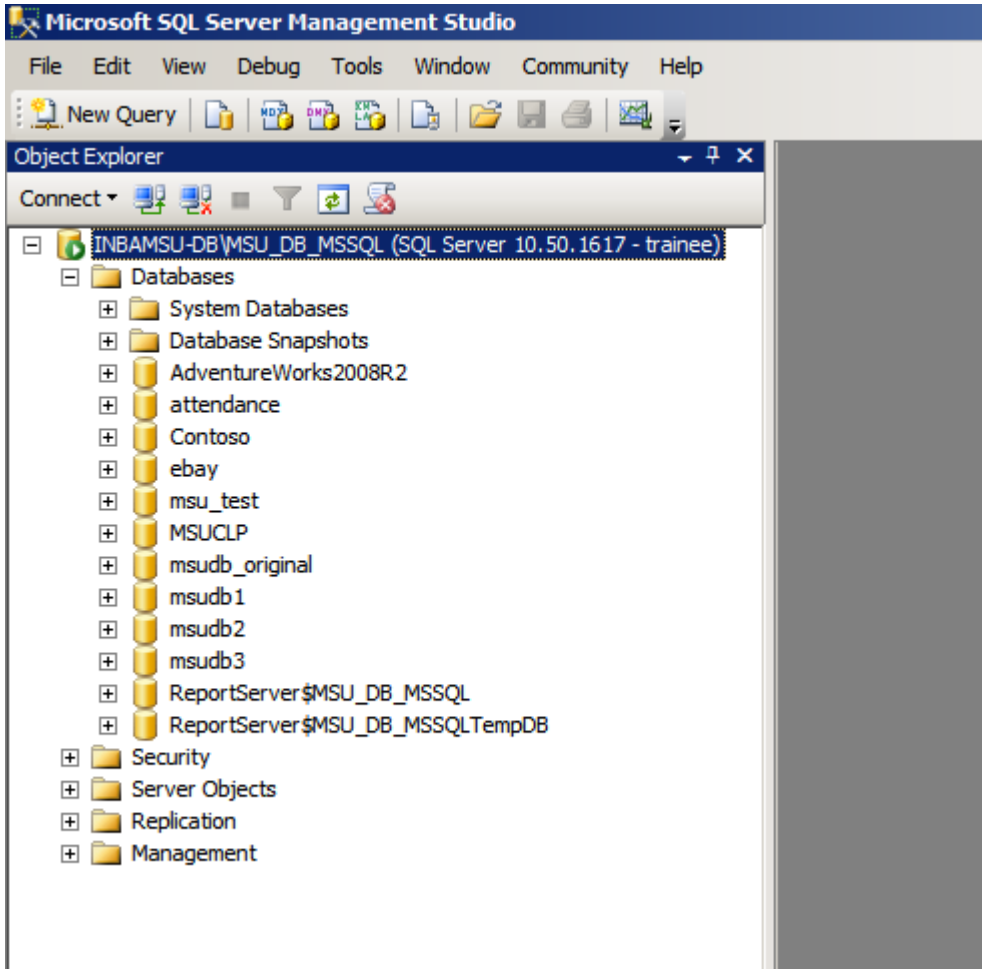

- Once you login to the server, the object explorer window on the left will list out all the objects
- Expand the **Database** object to display all the databases currently installed
- We will focus on the **Contoso** database for all our practice sessions

Open a query window by clicking on the New Query button

Open query window by clicking here

Select **Contoso** database from the drop down

Enter all your queries here

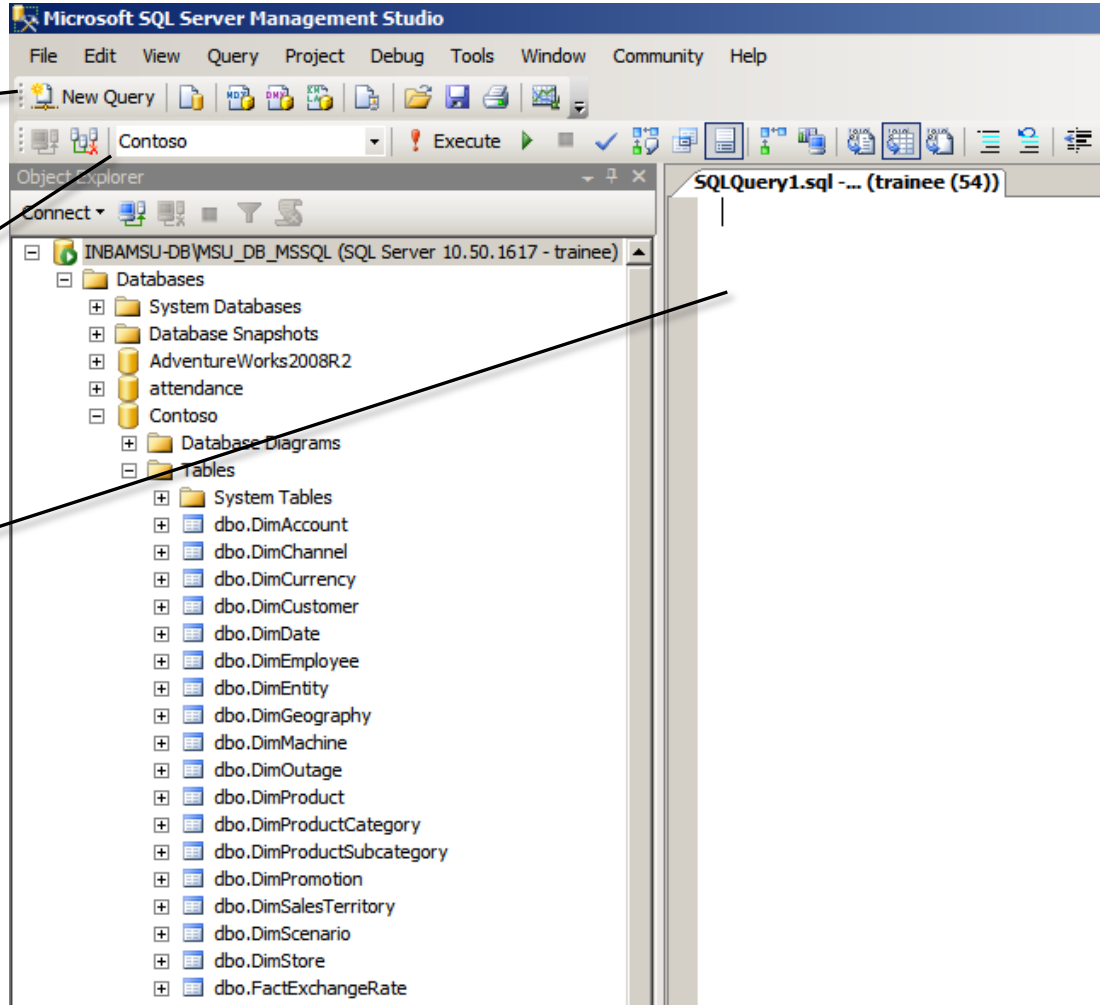

Your first SQL query

Enter the query and press **F5** or **Execute** to submit the query

Query **output**

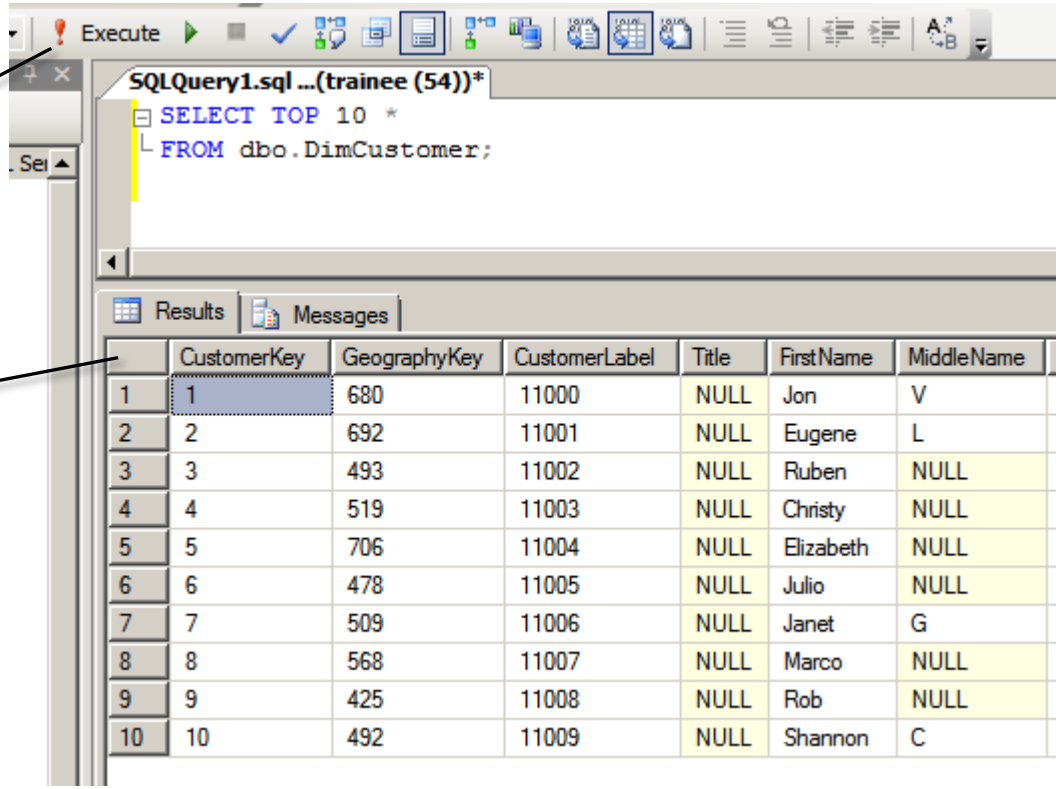

	CustomerKey	GeographyKey	CustomerLabel	Title	FirstName	MiddleName
1	1	680	11000	NULL	Jon	V
2	2	692	11001	NULL	Eugene	L
3	3	493	11002	NULL	Ruben	NULL
4	4	519	11003	NULL	Christy	NULL
5	5	706	11004	NULL	Elizabeth	NULL
6	6	478	11005	NULL	Julio	NULL
7	7	509	11006	NULL	Janet	G
8	8	568	11007	NULL	Marco	NULL
9	9	425	11008	NULL	Rob	NULL
10	10	492	11009	NULL	Shannon	C

The SELECT query extracts rows from a table

```
SELECT *
FROM dbo.DimProductCategory;
```

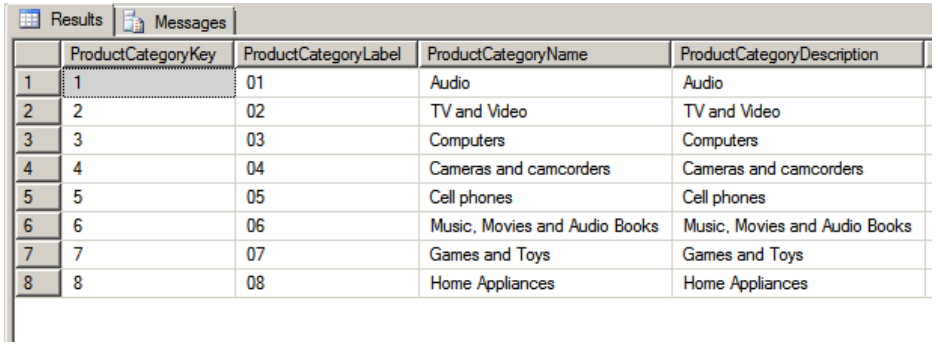

Results | Messages

	ProductCategoryKey	ProductCategoryLabel	ProductCategoryName	ProductCategoryDescription
1	1	01	Audio	Audio
2	2	02	TV and Video	TV and Video
3	3	03	Computers	Computers
4	4	04	Cameras and camcorders	Cameras and camcorders
5	5	05	Cell phones	Cell phones
6	6	06	Music, Movies and Audio Books	Music, Movies and Audio Books
7	7	07	Games and Toys	Games and Toys
8	8	08	Home Appliances	Home Appliances

- The **SELECT** query is used to select rows from a table in a database
- The asterisk (*) indicates all rows and all columns
- The **FROM** statement indicates the table you want to access

You can specify the column names to focus on only a few fields

```
SELECT
ProductCategoryLabel,
ProductCategoryName
FROM dbo.DimProductCategory;
```

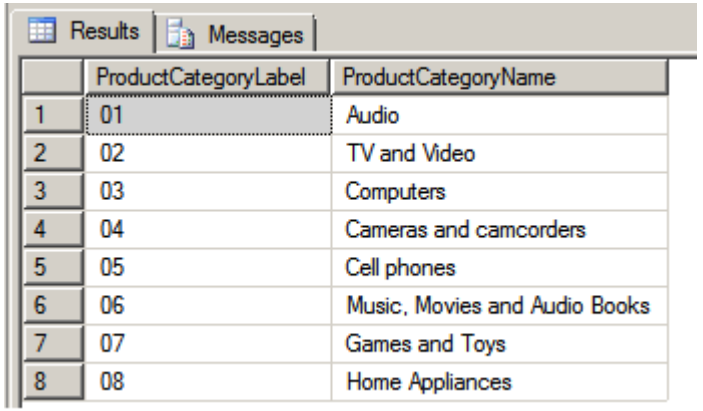

Results | Messages

	ProductCategoryLabel	ProductCategoryName
1	01	Audio
2	02	TV and Video
3	03	Computers
4	04	Cameras and camcorders
5	05	Cell phones
6	06	Music, Movies and Audio Books
7	07	Games and Toys
8	08	Home Appliances

- You can specify the columns you want to display by listing out the names separated by a comma
- SQL is not case sensitive, hence lower case syntax is also correct

Use the WHERE statement to apply conditions on rows

```
SELECT
        FirstName,
        LastName,
        YearlyIncome,
        Gender,
        Occupation
FROM dbo.DimCustomer
WHERE YearlyIncome > 100000;
```

Results | Messages

	FirstName	LastName	YearlyIncome	Gender	Occupation
1	Donald	Gonzalez	160000.00	M	Management
2	Damien	Chander	170000.00	M	Management
3	Savannah	Baker	120000.00	F	Management
4	Angela	Butler	130000.00	F	Management
5	Alyssa	Cox	130000.00	F	Management
6	Sarah	Thomas	110000.00	F	Management
7	Nicholas	Robinson	110000.00	M	Management
8	Jose	Flores	110000.00	M	Management
9	Molly	Rodriguez	120000.00	F	Management
10	April	Anand	160000.00	F	Management
11	Devin	Martin	170000.00	M	Management

- You can specify multiple conditions by separating each condition by the keywords **AND** or **OR**
- For better readability enclose the conditions with brackets

The ORDER statement sorts the output in ascending or descending order of a variable

```
SELECT
        FirstName,
        YearlyIncome,
        Occupation
FROM dbo.DimCustomer
WHERE YearlyIncome > 100000
ORDER BY YearlyIncome;
```

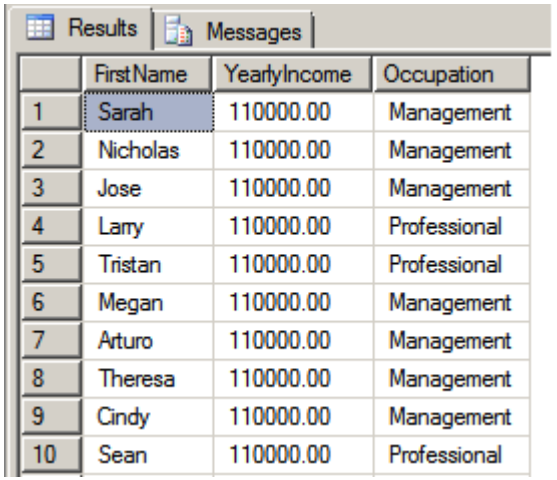

Results | Messages

	FirstName	YearlyIncome	Occupation
1	Sarah	110000.00	Management
2	Nicholas	110000.00	Management
3	Jose	110000.00	Management
4	Larry	110000.00	Professional
5	Tristan	110000.00	Professional
6	Megan	110000.00	Management
7	Arturo	110000.00	Management
8	Theresa	110000.00	Management
9	Cindy	110000.00	Management
10	Sean	110000.00	Professional

- The default sort order is ascending – to sort in descending order, use the keyword **DESC** after the variable name

The DISTINCT keyword identifies unique rows from a table

```
SELECT DISTINCT
        Occupation
FROM dbo.DimCustomer;
```

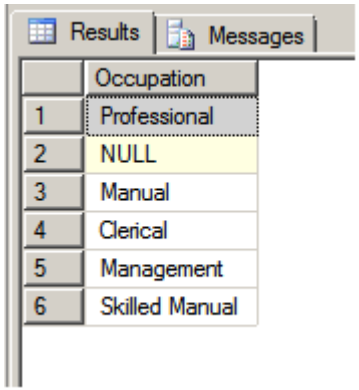

	Occupation
1	Professional
2	NULL
3	Manual
4	Clerical
5	Management
6	Skilled Manual

- Specifying multiple column names with **DISTINCT** keyword results in selecting unique combinations of all the columns from the table

The GROUP BY statement is used with summary functions to roll up tables

```
SELECT
        Occupation,
        Gender,
        COUNT(CustomerKey) AS
TotalCustomers
FROM dbo.DimCustomer
GROUP BY
        Occupation,
        Gender;
```

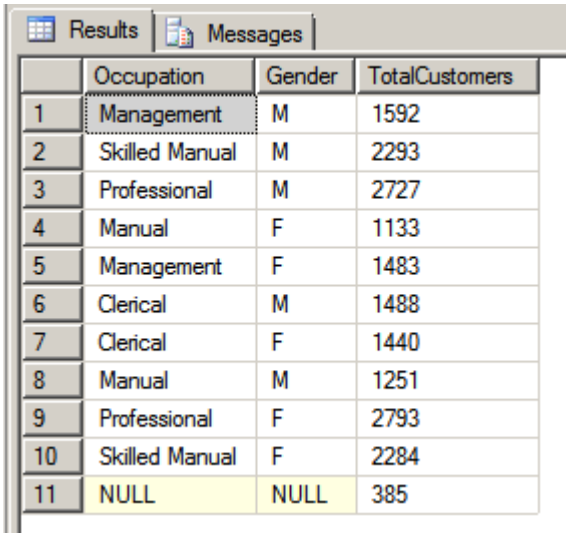

Results | Messages

	Occupation	Gender	TotalCustomers
1	Management	M	1592
2	Skilled Manual	M	2293
3	Professional	M	2727
4	Manual	F	1133
5	Management	F	1483
6	Clerical	M	1488
7	Clerical	F	1440
8	Manual	M	1251
9	Professional	F	2793
10	Skilled Manual	F	2284
11	NULL	NULL	385

- Various summary functions are available in SQL language like **SUM**, **AVG**, **MIN**, **MAX**, **STDEV**
- Additional functions may be available which can be specific to the database you are using
- A simple **COUNT(*)** will count the number of rows

The HAVING statement applies conditions on a summary variable

```
SELECT
        Occupation,
        Gender,
        COUNT(CustomerKey) AS
        TotalCustomers
FROM dbo.DimCustomer
GROUP BY
        Occupation, Gender
HAVING COUNT(CustomerKey) >
        2000
```

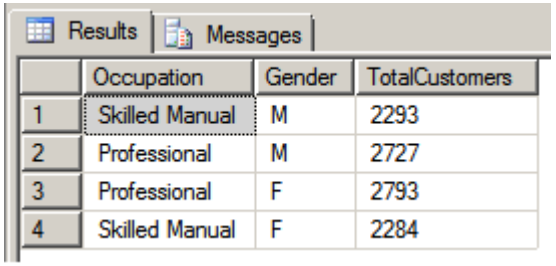

Results | Messages

	Occupation	Gender	TotalCustomers
1	Skilled Manual	M	2293
2	Professional	M	2727
3	Professional	F	2793
4	Skilled Manual	F	2284

- **HAVING** clause works on aggregation of data rather than the rows of the data prior to the aggregation

A **SUBQUERY** is a query on the data that is found within another query statement

```
SELECT
        FirstName,
        EmailAddress
FROM dbo.DimCustomer
WHERE GeographyKey IN (
        SELECT GeographyKey
        FROM dbo.DimGeography
        WHERE   CityName =
        'Beverley Hills');
```

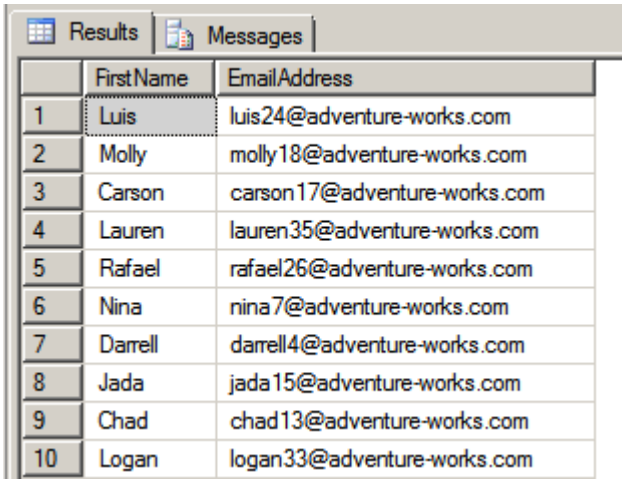

Results | Messages

	FirstName	EmailAddress
1	Luis	luis24@adventure-works.com
2	Molly	molly18@adventure-works.com
3	Carson	carson17@adventure-works.com
4	Lauren	lauren35@adventure-works.com
5	Rafael	rafael26@adventure-works.com
6	Nina	nina7@adventure-works.com
7	Darrell	darrell4@adventure-works.com
8	Jada	jada15@adventure-works.com
9	Chad	chad13@adventure-works.com
10	Logan	logan33@adventure-works.com

- The sub query can only return one column at a time and the inner query gets executed before the outer one

JOIN keyword is used to query data from two or more tables based on a relationships between the columns in these tables.

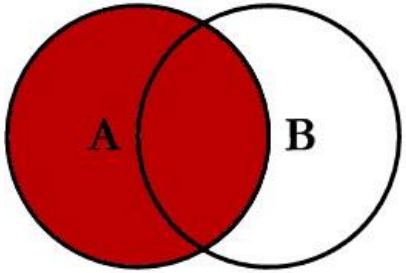


```
SELECT <select_list>
FROM TableA A
LEFT JOIN TableB B
ON A.Key = B.Key
```

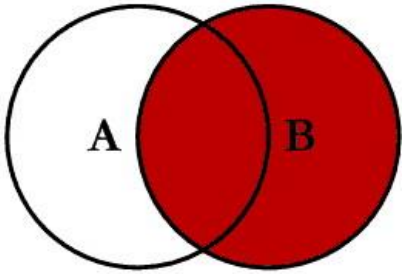


```
SELECT <select_list>
FROM TableA A
RIGHT JOIN TableB B
ON A.Key = B.Key
```

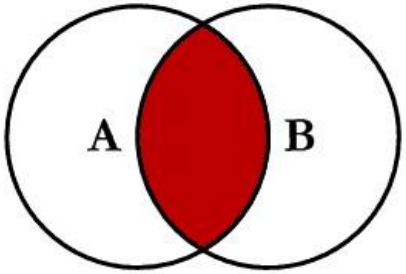


```
SELECT <select_list>
FROM TableA A
INNER JOIN TableB B
ON A.Key = B.Key
```

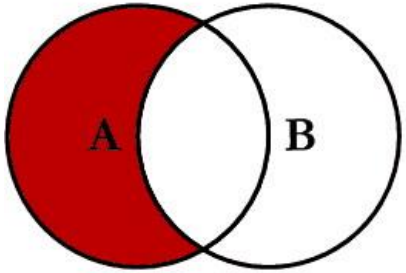


```
SELECT <select_list>
FROM TableA A
LEFT JOIN TableB B
ON A.Key = B.Key
WHERE B.Key IS NULL
```

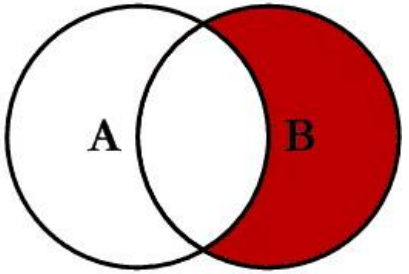


```
SELECT <select_list>
FROM TableA A
RIGHT JOIN TableB B
ON A.Key = B.Key
WHERE A.Key IS NULL
```

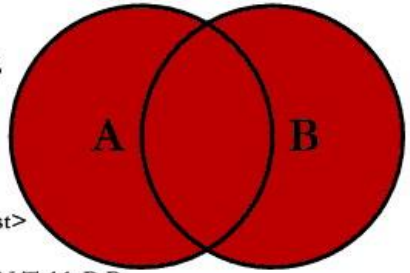


```
SELECT <select_list>
FROM TableA A
FULL OUTER JOIN TableB B
ON A.Key = B.Key
```

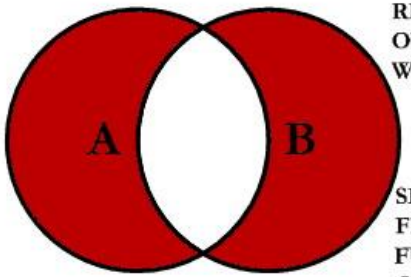


```
SELECT <select_list>
FROM TableA A
FULL OUTER JOIN TableB B
ON A.Key = B.Key
WHERE A.Key IS NULL
OR B.Key IS NULL
```

An SQL join can merge multiple tables at a time

```
SELECT
        t1.FirstName,
        t2.StateProvinceName,
        t2.CityName
FROM dbo.DimCustomer AS t1
INNER JOIN dbo.DimGeography
AS t2 ON (t1.GeographyKey =
t2.GeographyKey)
```

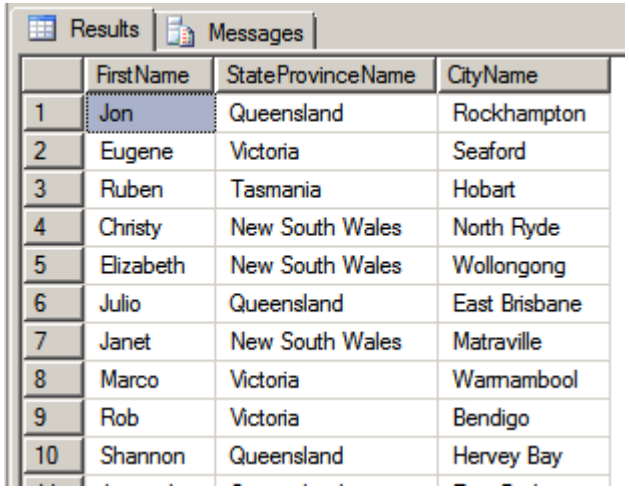

Results | Messages

	FirstName	StateProvinceName	CityName
1	Jon	Queensland	Rockhampton
2	Eugene	Victoria	Seaford
3	Ruben	Tasmania	Hobart
4	Christy	New South Wales	North Ryde
5	Elizabeth	New South Wales	Wollongong
6	Julio	Queensland	East Brisbane
7	Janet	New South Wales	Matraville
8	Marco	Victoria	Warrnambool
9	Rob	Victoria	Bendigo
10	Shannon	Queensland	Hervey Bay

- The advantage of using table alias (t1, t2 etc.) is that merging column can be of different names in different tables
- Joins can happen on multiple columns too

Further reading and practice

- Data warehousing concepts (Link)
- Interactive SQL Tutorial (Link)
- Beginning SQL Server 2008 for Developers (Chapters 11 and 12) (Link)
- SQL Server 2012: T-SQL Fundamentals (Link)
- SQL Server 2008: T-SQL Querying (Link)